CYBER DECEPTION AI

AI-Based Cyber Deception System for Intelligent Honeypot Placement and Attacker Path Prediction

ARTIFICIAL INTELLIGENCE – PROJECT

Submitted in partial fulfillment of the requirements of the subject

Artificial Intelligence

in

Department of Cybersecurity and Intelligent Systems

by

CHIDWILAS ROUTHU

24BYB0092



VIT®
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

School of Computer Science Engineering and Information Systems [SCORE]

VIT, Vellore

October, 2026

DECLARATION

I hereby declare that the Artificial Intelligence project entitled **"Cyber Deception AI"** submitted by me to the School of Computer Science Engineering and Information Systems, VIT, Vellore, is a record of bonafide work carried out by me.

I further declare that the work reported in this project has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university.

Place: Vellore

Date:

Signature of the Candidate
Chidwilas Routhu (24BYB0092)

CERTIFICATE

This is to certify that the Artificial Intelligence project entitled **"Cyber Deception AI"** submitted by **Chidwilas Routhu (24BYB0092)**, SCORE, VIT, Vellore, is a record of work carried out as per the VIT code of academic and research ethics.

The contents of this report have not been submitted and will not be submitted either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university. The project fulfils the requirements and regulations of the University and in my opinion meets the necessary standards for submission.

Place: Vellore

Date:

Signature of the VIT-SCORE Guide

CONTENTS

		Page No.
1	**TITLE**	i
	Declaration	ii
	Certificate	iii
	List of Figures	v
	List of Tables	vi
	List of Abbreviations	vii
2	**ABSTRACT**	1
3	**INTERACTION**	2
4	**LITERATURE SURVEY**	5
5	**SYSTEM ARCHITECTURE**	7
6	**MODULE DESCRIPTION**	10
7	**RESULTS AND IMPLEMENTATIONS**	13
8	**CONCLUSION**	22
9	**REFERENCES**	23
10	**DEPLOYMENT**	25

LIST OF FIGURES

Figure No.	Title	Page No.
3.1	Dashboard After an Analysis Run	2
3.2	Network Topology Page With Asset Inspector	3
5.1	System Architecture	7
5.2	Workflow of One Simulation Request	8
7.1	Attack Path Analysis Workspace	13
7.2	Simulation Result and Deception Scoring	14
7.3	Greedy Placement Candidate Scores	15
7.4	Greedy vs Hill Climbing vs Genetic Algorithm	16
7.5	Experiment Lab Strategy Comparison	17
7.6	Advanced Analysis and Defense Planning	17
7.7	Dashboard Before the First Analysis	19
7.8	Stepped Simulation	19
7.9	Saved Scenarios Page	20
7.10	Methodology Page	20
10.1	Live Deployment on GitHub Pages	25

LIST OF TABLES

Table No.	Title	Page No.
4.1	Summary of Surveyed Literature	6
5.1	Built-In Simulated Network	8
7.1	Search Algorithms on the Built-In Network	13
7.2	Greedy Deception Scores	14
7.3	Comparison of the Three Placement Algorithms	15
7.4	Interception Rate by Strategy	16
7.5	Functional Test Cases	18
7.6	Automated Test Results	18
10.1	Deployment Summary	25

LIST OF ABBREVIATIONS

A*	A-star informed search
AI	Artificial Intelligence
API	Application Programming Interface
BFS	Breadth-First Search
CORS	Cross-Origin Resource Sharing
DFS	Depth-First Search
GA	Genetic Algorithm
HC	Hill Climbing
JSON	JavaScript Object Notation
ORM	Object-Relational Mapping
REST	Representational State Transfer
STRIPS	Stanford Research Institute Problem Solver
UCS	Uniform-Cost Search
UI	User Interface

2. ABSTRACT

Honeypots and other deception resources give defenders high-confidence alerts, because no legitimate user should ever interact with them. An organisation cannot place a decoy on every host, however, so the value of deception depends on *where* a limited number of decoys are placed. **Cyber Deception AI** is a simulation and decision-support system that addresses this placement problem on a modelled network.

The network is represented as a weighted, undirected graph in which nodes are hosts and edge weights represent a simulated traversal difficulty. An attacker agent predicts its route from an entry point to a critical asset using four hand-implemented search algorithms: Breadth-First Search, Depth-First Search, Uniform-Cost Search and A*. A defender agent then selects honeypot locations within a fixed budget using three approaches: a transparent **Greedy** deception score, **Hill Climbing** with random restarts, and a **Genetic Algorithm**. All three strategies are evaluated with the same placement-fitness function.

The system is implemented as a full-stack application with a React/TypeScript front end and a FastAPI/NetworkX back end. Saved scenarios are persisted with SQLAlchemy in PostgreSQL. It is publicly deployed on free services: GitHub Pages hosts the front end, and Render hosts the API and database. On the built-in seven-host network, the A* route Internet → Web Server → Application Server → Database has cost 8. With a budget of two honeypots, all three algorithms reach the same fitness (6.2667) and cover all three cheapest attack routes. In a 20-run experiment with one honeypot, random placement intercepts 50% of attacks, while Greedy, Hill Climbing and the Genetic Algorithm intercept 100%.

The experiments are **simulated**, and the network contains a chokepoint: every route passes through the Web Server. The results therefore demonstrate a working, explainable pipeline rather than real-world detection effectiveness. The system uses classical AI techniques (search, heuristics, optimisation, adversarial search and Bayesian reasoning) and contains no trained machine-learning model.

Keywords: cyber deception, honeypot placement, attack paths, graph search, greedy optimisation, hill climbing, genetic algorithm.

3. INTERACTION

This chapter describes how a user interacts with Cyber Deception AI: the inputs supplied, the processing that each action triggers, and the outputs shown. All behaviour described here was exercised on the running application, both locally and on the public deployment described in Section 10.

3.1 Opening the Application

The user opens the deployed site at https://sid07-pro.github.io/cyber-deception-ai/ or, for local use, http://localhost:5173. A left sidebar groups the pages into four sections:

- *Overview:* Dashboard, Network Topology.
- *Analysis:* Attack Path Analysis, Honeypot Placement, Advanced Analysis, Defense Planning.
- *Evaluation:* Experiment Lab, Saved Scenarios.
- *Reference:* Methodology.

The header shows a “Simulated environment” badge. A status indicator polls the real `/api/health` endpoint every 30 seconds and reports the API version and database status (Fig. 3.1).

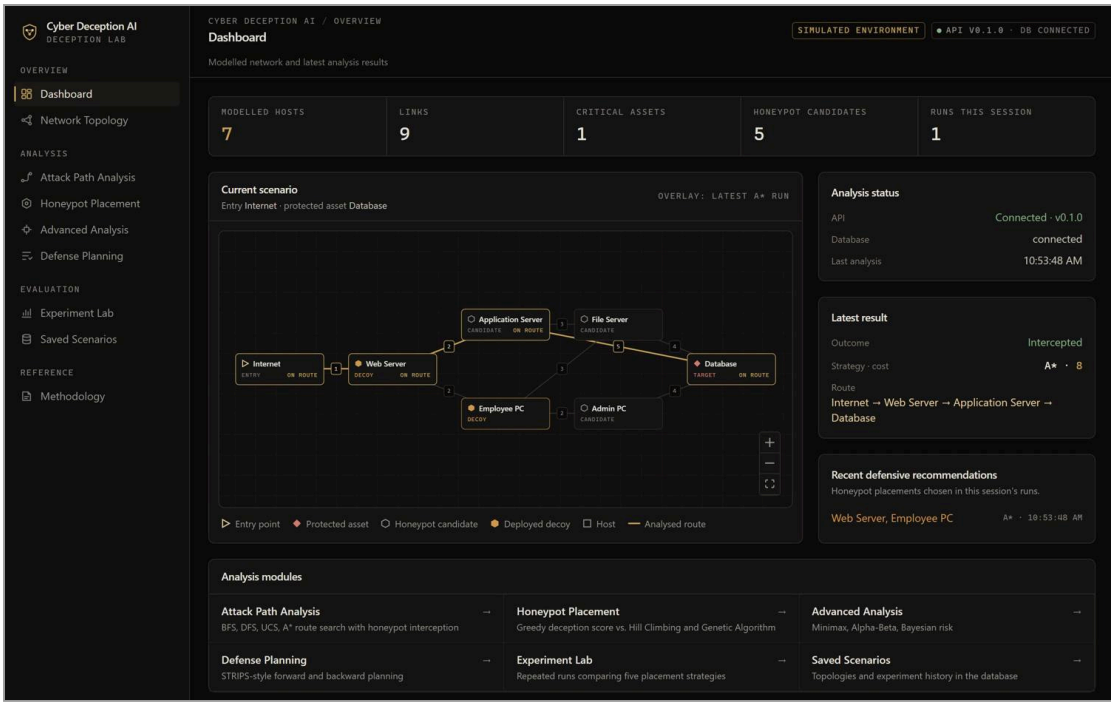


Fig. 3.1 Dashboard After an Analysis Run

3.2 Supplying Network Information

The built-in network is loaded automatically from `GET /api/network`. On the Network Topology page (Fig. 3.2), the user can click any host in the graph or the host table to inspect its type, criticality, degree, deployment cost and links, and can set it as the analysis start or target. Custom topologies can be stored as saved scenarios through the API and simulated from the Saved Scenarios page. There is no graphical topology editor.

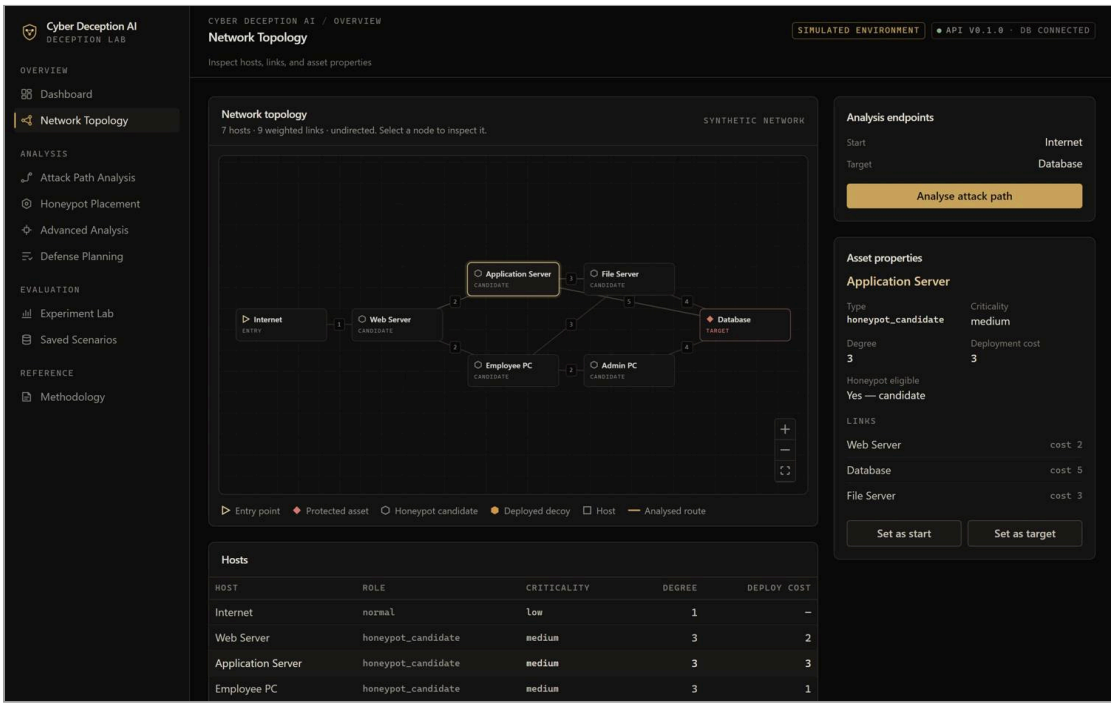


Fig. 3.2 Network Topology Page With Asset Inspector

3.3 Selecting and Executing Algorithms

- **Attack Path Analysis:** the user selects a search strategy (BFS, DFS, UCS or A*), a start node, a target node and a honeypot budget (0–5), then presses *Run Simulation*. A stepped mode moves the attacker one hop at a time.
- **Honeypot Placement:** *Evaluate placements* runs the Greedy score, Hill Climbing and the Genetic Algorithm on the same scenario and budget.
- **Advanced Analysis:** runs Minimax, Alpha-Beta, Bayesian risk, Hill Climbing or the Genetic Algorithm individually.
- **Defense Planning:** generates a forward or backward state-space plan.
- **Experiment Lab:** repeats runs and compares five placement strategies. Results can be exported as JSON or CSV, or saved to the database.

3.4 Presentation of Outputs and Error Handling

Results show the following:

- the predicted route, drawn in gold on the graph;
- the route cost and the number of nodes explored;
- the deployed honeypots, tagged “Decoy”;
- the point where the attacker was intercepted;
- the full deception-score table;
- the three cheapest alternative routes;
- a plain-language explanation of each decision.

Invalid input is rejected at two levels. In the interface, the *Run* and *Evaluate* buttons are disabled when start and target are the same node. On the back end, every request is validated:

- an invalid request body returns HTTP 422;
- an unknown node or scenario returns 404;
- an unavailable database returns 503.

The interface displays the server’s message in an error panel.

4. LITERATURE SURVEY

This chapter reviews research on honeypots and cyber deception, game-theoretic defence, graph-based security analysis, search, and the three optimisation approaches compared in this project.

4.1 Honeypots and Cyber Deception

Spitzner [1] established the honeypot as a security resource whose value lies in being probed or attacked. He distinguished low-interaction honeypots, which emulate services, from high-interaction honeypots, which are real monitored systems. Provos [2] introduced Honeyd, a framework that simulates many virtual hosts with configurable network personalities and shows that large decoy networks can be emulated cheaply. Both works focus on building and operating decoys rather than deciding where a limited number should be placed, which is the focus of this project. MITRE Engage [3] frames deception as part of adversary engagement, aligned with defender goals.

4.2 Game-Theoretic and Strategic Deception

Pawlick, Colbert and Zhu [4] survey game-theoretic models of defensive deception and classify them into six categories: perturbation, moving-target defence, obfuscation, mixing, honey-x and attacker engagement. Tambe [5] describes security games in which a defender with limited resources commits to a strategy and an attacker best-responds. This project borrows the attacker–defender framing but implements depth-limited adversarial search (Minimax and Alpha-Beta [6]) with a hand-designed utility. It does not compute a game-theoretic equilibrium.

4.3 Graph-Based Network Security Analysis

Sheyner et al. [7] showed that attack graphs can be generated automatically from a network model and used to find the paths that reach a goal. Ou, Govindavajhala and Appel [8] proposed MulVAL, a logic-based analyser that derives attack graphs from vulnerability and configuration data. Durkota et al. [9] combined attack graphs with game theory to compute defensive hardening, including honeypot deployment. Cyber Deception AI uses a simpler weighted host graph without vulnerability data. Its contribution is to connect route prediction directly to decoy placement in an explainable, interactive tool.

4.4 Search Algorithms

Dijkstra’s algorithm [10] computes minimum-cost paths in graphs with non-negative weights; Uniform-Cost Search is its single-target form. Hart, Nilsson and

Raphael [11] introduced A*, which orders expansion by f(n) = g(n) + h(n) and is optimal when the heuristic h is admissible. Yen [12] gave an algorithm for the K shortest loopless paths, which this project uses through NetworkX [13] to enumerate alternative attack routes. Standard treatments appear in Cormen et al. [14] and Russell and Norvig [15].

4.5 Greedy Optimisation, Hill Climbing and Genetic Algorithms

Nemhauser, Wolsey and Fisher [16] proved that greedy selection achieves a (1 − 1/e) approximation when maximising a monotone submodular set function, such as coverage, under a cardinality constraint. That guarantee applies to marginal-gain greedy selection on a submodular objective. This project's greedy method instead ranks nodes by an independent per-node score, so the guarantee does not transfer. Hill Climbing is a local-search method that is prone to local optima, a weakness that random restarts mitigate [15]. Genetic Algorithms, introduced by Holland [17] and popularised by Goldberg [18], evolve a population of solutions through selection, crossover and mutation. STRIPS [19] provides the action representation used by the planning module.

Table 4.1 Summary of Surveyed Literature

Ref.	Problem / approach	Relevance and difference
[1], [2]	Honeypot design; virtual honeypot framework	Defines the decoy concept; does not address budgeted placement
[4], [5]	Game-theoretic deception; security games	Attacker–defender framing; this project uses adversarial search, not equilibria
[7]–[9]	Attack graphs; logic-based analysis; attack-graph games	Route modelling; this project uses a simpler weighted host graph
[10]–[12]	Shortest paths; A*; K shortest loopless paths	Directly implemented for attacker routes and alternatives
[16]–[18]	Greedy approximation; genetic algorithms	Basis of the placement methods; [16]'s guarantee does not apply to the per-node score

5. SYSTEM ARCHITECTURE

The application is a three-tier web system (Fig. 5.1):

- a React single-page application calls a FastAPI REST service over JSON;
- the service layer composes framework-independent algorithm modules written in Python with NetworkX;
- SQLAlchemy persists saved scenarios and experiment results.

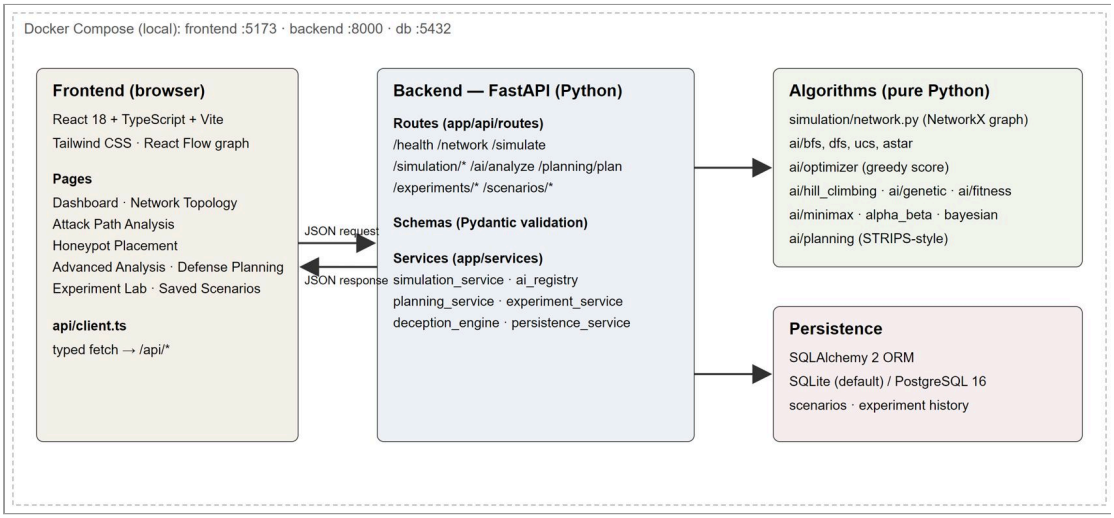


Fig. 5.1 System Architecture

5.1 Components

- **Front end** (`frontend/src`): React 18, TypeScript, Vite 5, Tailwind CSS and React Flow 11. A typed client issues all API requests.
- **API layer** (`backend/app/api/routes`): handles routing, status codes and error mapping. Pydantic schemas validate every request.
- **Service layer** (`backend/app/services`): the simulation, AI registry, planning, experiment, deception-engine and persistence services. They contain no HTTP code.
- **Algorithm layer** (`backend/ai`, `backend/simulation`): graph model, search, placement, game search, Bayesian reasoning and planning.
- **Persistence**: SQLAlchemy 2 ORM. It uses SQLite by default and PostgreSQL under Docker and on Render.

5.2 Graph Representation

The built-in network (`backend/simulation/network.py`) is an undirected NetworkX graph G = (V, E) with |V| = 7 and |E| = 9:

- each node carries a role, and each edge carries a weight;

- Internet is the entry point and Database is the critical target;
- the other five hosts are honeypot candidates.

All weights are hand-chosen simulation values (Table 5.1). Internet has a single neighbour, Web Server, so every attack route passes through Web Server.

Table 5.1 Built-In Simulated Network

Edge	Weight	Edge	Weight
Internet – Web Server	1	Employee PC – Admin PC	2
Web Server – Application Server	2	Admin PC – Database	4
Web Server – Employee PC	2	File Server – Database	4
Application Server – Database	5	Employee PC – File Server	3
Application Server – File Server	3		

5.3 Data Flow

Fig. 5.2 shows how one simulation request is processed. The request moves from user input through validation, graph construction, attacker search, defender placement and evaluation, and finally to the response rendered by the interface.

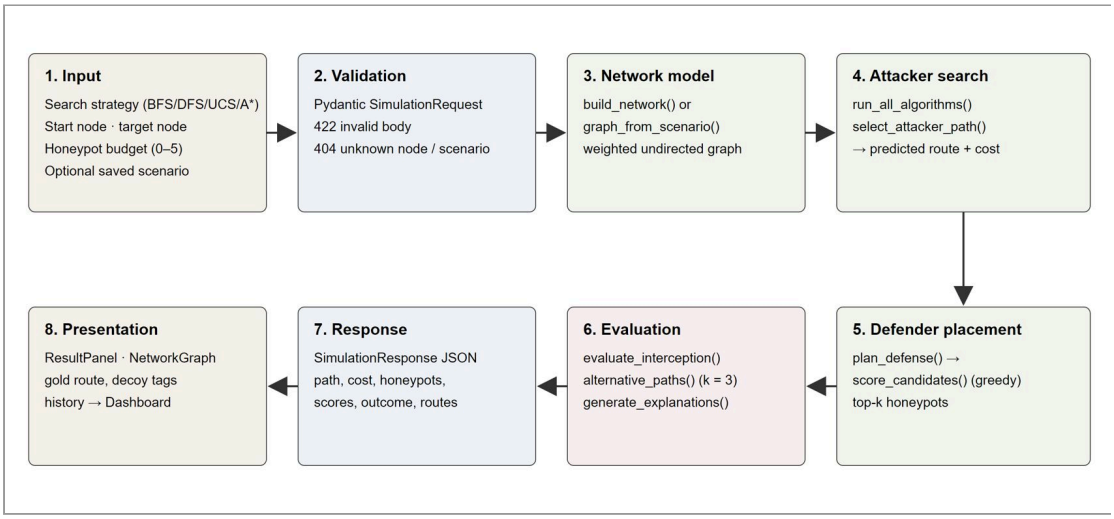


Fig. 5.2 Workflow of One Simulation Request

5.4 Deployment Architecture

The production deployment has three parts:

- the static front end, built by a GitHub Actions workflow and served by GitHub Pages;
- the FastAPI service, run by Render as a free web service;
- a free Render PostgreSQL database for persistence.

Locally, Docker Compose runs the same components as three services: `db`, `backend` and `frontend`.

6. MODULE DESCRIPTION

This chapter describes each implemented module, giving its purpose, inputs, processing, outputs and source files.

6.1 Network Model

`build_network()` adds nodes with roles and weighted edges. `graph_from_scenario()` rebuilds a graph from stored nodes and edges and reads optional per-node deployment costs. Saved topologies are validated on save: node IDs must be unique, edge endpoints must exist, weights must be positive, and the start and target must be distinct nodes.

6.2 Attack-Route Prediction

Four hand-written searches share the interface `search(graph, start, goal) → (path, cost, nodes_explored)`:

- **BFS:** uses a FIFO queue and finds the fewest hops. Time O(V + E).
- **DFS:** uses an explicit stack and gives no optimality guarantee. Time O(V + E).
- **UCS:** uses a min-priority queue on g(n) and is optimal for positive weights. Time O((V + E) log V).
- **A*:** uses a priority queue on f(n) = g(n) + h(n), where h(n) is the hop distance to the target multiplied by the minimum edge weight. This heuristic is admissible, so A* is optimal.

`run_all_algorithms()` runs all four searches and supplies their routes to the defender; `select_attacker_path()` then commits to the chosen strategy's route. Alternative routes are the three cheapest loop-free paths from Yen's algorithm [12], computed by `alternative_paths()`.

6.3 Greedy Placement

The greedy method (`backend/ai/optimizer.py: score_candidates`) scores each candidate n independently, excluding the start and target, and selects the top k by Eqn. 6.1:

$$score(n) = [\, freq(n) + 1 / (hops(n, target) + 1) + deg(n) \,] / cost(n) \quad (6.1)$$

The terms in Eqn. 6.1 are defined as follows:

- **freq(n):** the number of candidate routes that contain n;
- **hops(n, target):** the unweighted distance from n to the target;
- **deg(n):** the node degree;

- **cost(n):** the simulated deployment cost (Web Server 2, Application Server 3, Employee PC 1, Admin PC 2, File Server 2).

The method is fast and transparent, but it ignores interactions between the selected decoys.

6.4 Placement Fitness

`backend/ai/fitness.py: evaluate()` scores a whole placement P using Eqn. 6.2:

$$F(P) = 4 \cdot coverage + 2 \cdot proximity + 1 \cdot criticality + 1.5 \cdot deception - 1 \cdot cost_{norm} \quad (6.2)$$

The terms in Eqn. 6.2 are defined as follows:

- **coverage:** the fraction of candidate routes that contain at least one node of P;
- **proximity:** the mean of 1/(1 + hops to target);
- **criticality:** the mean role weight;
- **deception:** the mean of degree divided by maximum degree;
- **$cost_{norm}$:** the normalised deployment cost.

Hill Climbing and the Genetic Algorithm maximise F, and the experiment service reports F for every strategy, so all strategies are compared on one objective.

6.5 Hill Climbing

`backend/ai/hill_climbing.py` performs hill climbing over placements of at most k eligible nodes. A placement's neighbours are those one swap, drop or add away. There are four restarts: the first starts from the highest-degree hosts and the others from random placements produced by a seeded generator. Each pass scans the neighbours in a deterministic order and accepts any improving neighbour as soon as it is found. The search stops when a full pass brings no improvement.

```
best ← None
for restart in 1..4:
    P ← highest-degree hosts (first restart) or random placement
    repeat:
        improved ← false
        for N in neighbours(P):            # swap, add or drop one
node
            if F(N) > F(P): P ← N; improved ← true
    until not improved
    best ← better of (best, P)
return best
```

6.6 Genetic Algorithm

`backend/ai/genetic.py` represents each individual as a set of exactly k eligible nodes. Its parameters are:

- population 12, evolved for 15 generations;
- tournament selection;
- crossover that mixes the two parents' nodes;
- mutation rate 0.2;
- a repair step that restores exactly k distinct nodes.

All randomness is seeded, so runs are reproducible.

6.7 Interception and Evaluation

`evaluate_interception()` returns the first honeypot on the attacker path and its position. `experiment_service.run_experiment()` repeats runs for five strategies (none, random, greedy, hill climbing and genetic algorithm). It aggregates the interception rate, attack success rate, path cost, steps, fitness and execution time.

6.8 Supporting Modules

- **Adversarial search** (`deception_game.py, minimax.py, alpha_beta.py`): MAX (the attacker) moves one hop; MIN (the defender) places a honeypot or passes. The terminal utility is ±50.
- **Bayesian risk** (`bayesian.py`): sequential Bayes updates using assumed likelihoods.
- **Planning** (`planning.py`): STRIPS-style forward search and backward regression, both breadth-first over states.
- **Back-end API and persistence**: health, network, simulation, session, AI, planning, experiment and scenario endpoints, with SQLAlchemy models for scenarios and experiment records.

7. RESULTS AND IMPLEMENTATIONS

All results were produced by the project's own code on the built-in **simulated** network. The screenshots are genuine captures of the running application. Runtimes are machine-dependent and are reported only as an indication.

7.1 Attack-Route Prediction

Table 7.1 Search Algorithms on the Built-In Network

Algorithm	Predicted route (Internet → Database)	Cost	Nodes explored
BFS	Internet → Web Server → Application Server → Database	8	5
DFS	same	8	4
UCS	same	8	7
A*	same	8	7

All four algorithms return the same route (Table 7.1). A* explores as many nodes as UCS because its admissible heuristic is weak on this graph. The three cheapest loop-free routes cost 8, 9 and 10 (Fig. 7.1, Fig. 7.2).

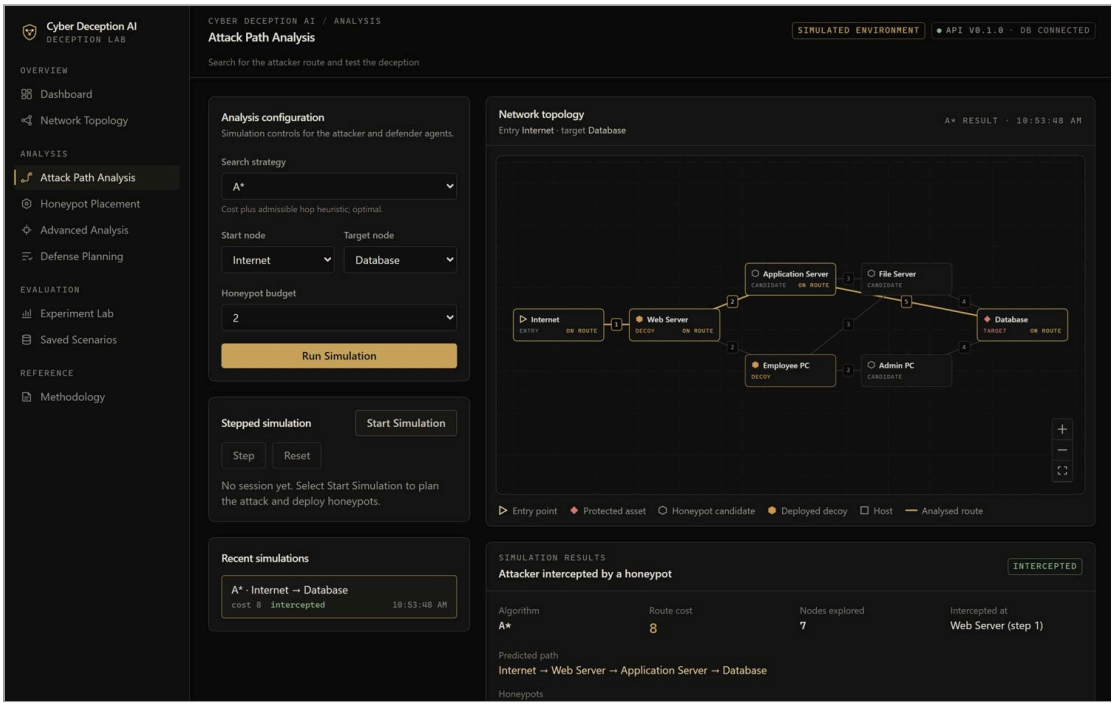


Fig. 7.1 Attack Path Analysis Workspace

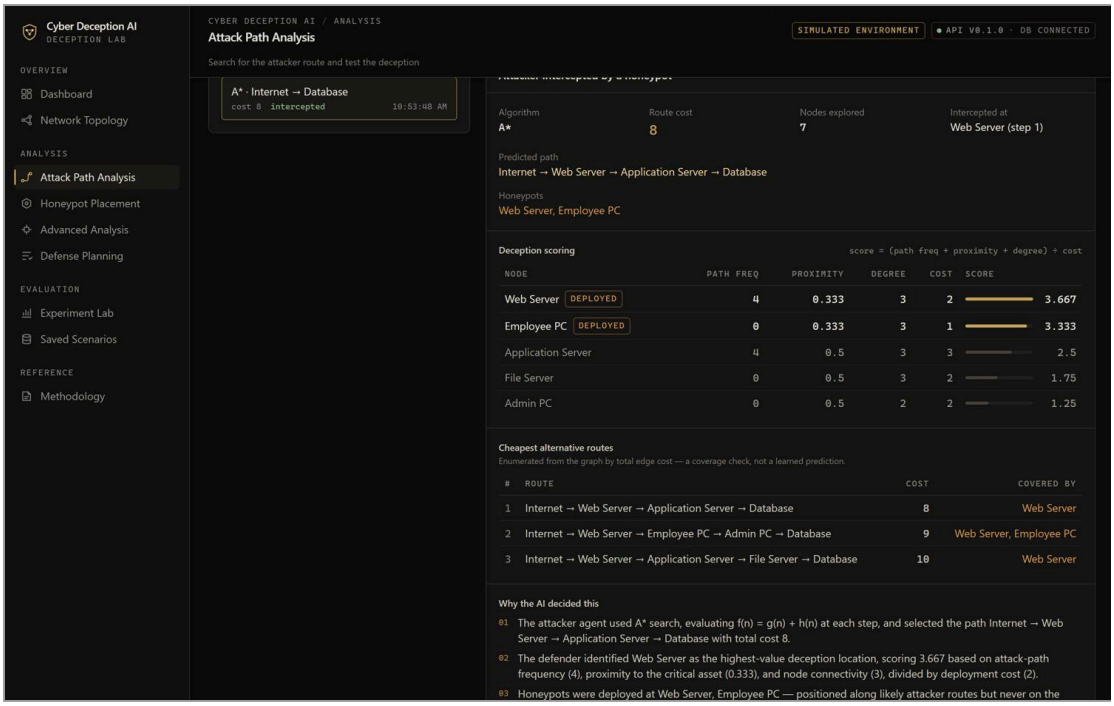


Fig. 7.2 Simulation Result and Deception Scoring

7.2 Greedy Placement

Table 7.2 Greedy Deception Scores (A*, Budget k = 2)

Candidate	freq	proximity	degree	cost	score	Selected
Web Server	4	0.333	3	2	3.667	Yes
Employee PC	0	0.333	3	1	3.333	Yes
Application Server	4	0.5	3	3	2.5	No
File Server	0	0.5	3	2	1.75	No
Admin PC	0	0.5	2	2	1.25	No

By Eqn. 6.1, score(Web Server) = (4 + 1/3 + 3) / 2 = 3.667. Employee PC is selected mainly because of its low cost and degree, even though no predicted route uses it (Fig. 7.3).

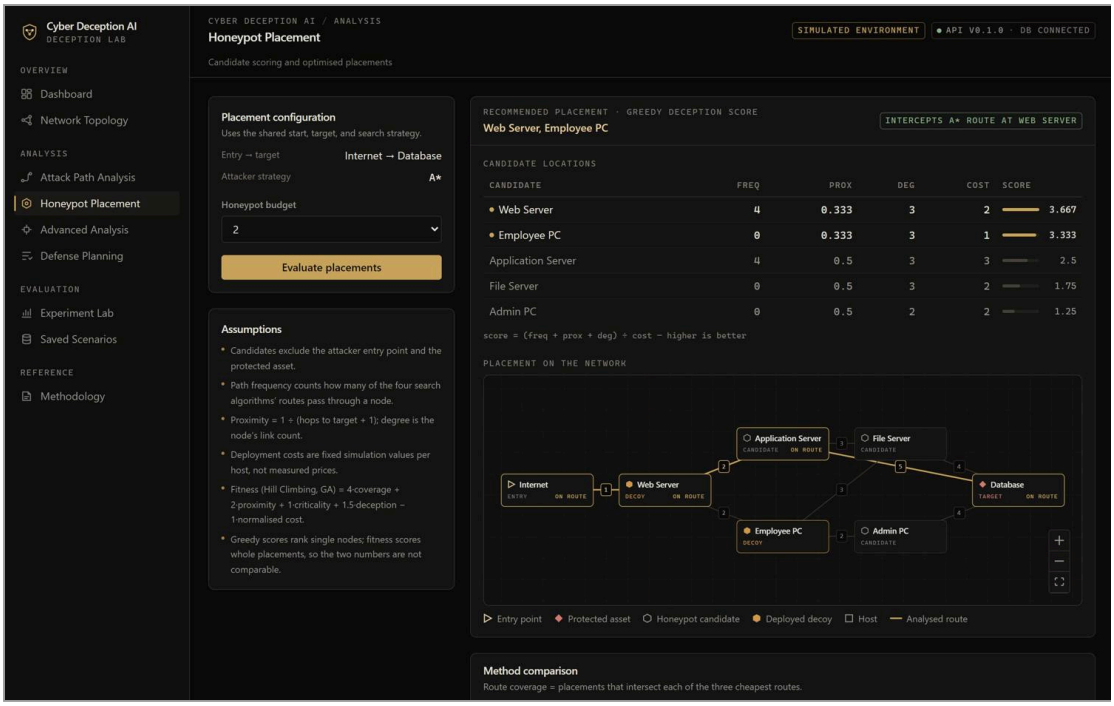


Fig. 7.3 Greedy Placement Candidate Scores

7.3 Comparison of Greedy, Hill Climbing and Genetic Algorithm

The three algorithms were run on the same network, attacker strategy (A*), candidate set and budget (k = 2). Routes covered counts how many of the three cheapest routes intersect the placement. Fitness is given by Eqn. 6.2, and runtime is the average per run over 20 runs.

Table 7.3 Comparison of the Three Placement Algorithms

Algorithm	Placement selected	Routes covered	Coverage	Fitness	Evaluations	Avg. runtime
Greedy	Web Server, Employee PC	3 / 3	100%	6.2667	—	0.13 ms
Hill Climbing	Application Server, Employee PC	3 / 3	100%	6.2667	44	2.28 ms
Genetic Algorithm	File Server, Web Server	3 / 3	100%	6.2667	194	13.16 ms

All three methods reach the same fitness with different placements, which shows that the objective has ties on this small network (Table 7.3, Fig. 7.4). Greedy is the fastest, but its advantage here comes from the chokepoint rather than from a general guarantee.

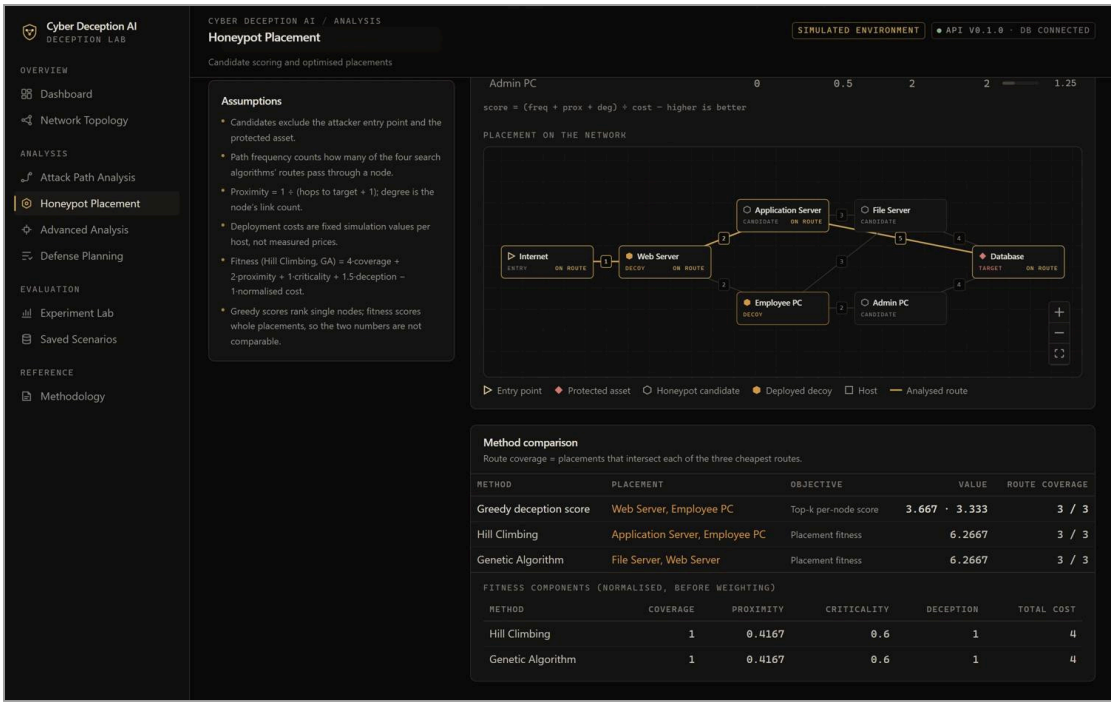


Fig. 7.4 Greedy vs Hill Climbing vs Genetic Algorithm

7.4 Interception Experiment

Table 7.4 Interception Rate by Strategy (A*, 20 Runs, Seed 0)

Strategy	Budget 1	Budget 2	Avg. fitness (budget 1)
No deception	0%	0%	0.0
Random placement	50%	85%	4.1917
Greedy	100%	100%	6.1
Hill Climbing	100%	100%	6.1
Genetic Algorithm	100%	100%	6.1

The attacker is deterministic, so only the random strategy varies between runs. The 100% interception rate follows from the chokepoint at Web Server and is not a real-world detection rate (Table 7.4, Fig. 7.5).

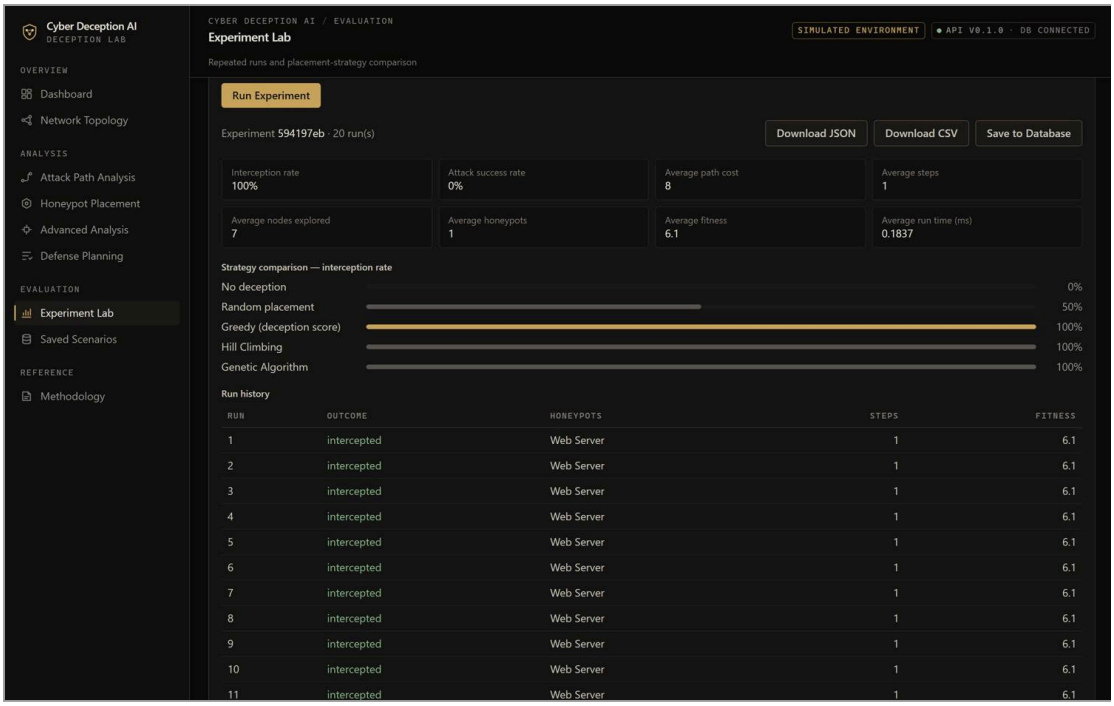


Fig. 7.5 Experiment Lab Strategy Comparison

7.5 Supporting Analyses

At depth 4, Minimax visited 69 states and Alpha-Beta 35 (14 branches pruned); both chose “move Web Server” with value 2.6667. The Bayesian module gives the Database a posterior of 0.969 from a prior of 0.70, based on assumed likelihoods. Forward planning (48 states explored) and backward regression (5 states) both found the same three-action plan, with cost 7 (Fig. 7.6).

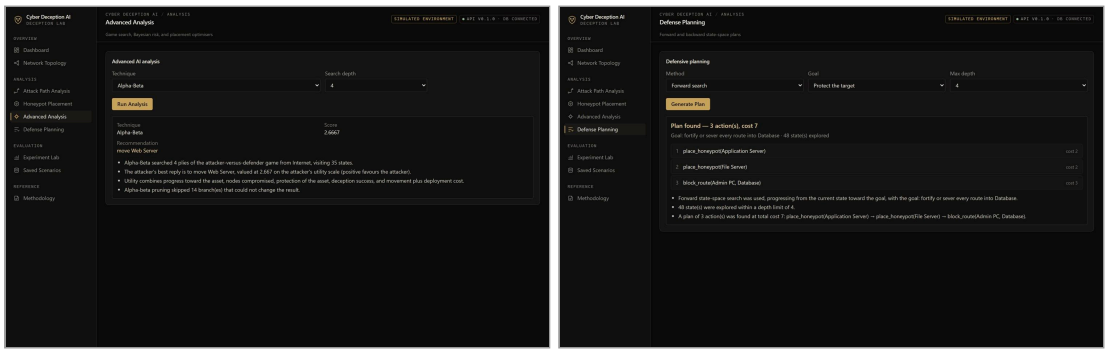


Fig. 7.6 Advanced Analysis (Alpha-Beta) and Defense Planning

7.6 Testing

Table 7.5 Functional Test Cases

#	Input	Expected behaviour	Actual result	Status
1	POST /api/simulate {algorithm: "UCS"}	Cheapest route, honeypots placed	Cost 8; Web Server, Employee PC; intercepted	Pass
2	honeypot_count = 99	Validation error	HTTP 422	Pass
3	start_node = "Mainframe"	Unknown-node error	HTTP 404 with known nodes	Pass
4	Disconnected target (test graph)	No path	Empty path, outcome no_path	Pass
5	Simulate saved scenario	Uses stored topology	HTTP 200, scenario_id returned	Pass
6	scenario_id = 999	Not found	HTTP 404	Pass
7	Stepped session, one step	Attacker moves one hop	At Web Server, intercepted, step 1	Pass
8	Start = target in UI	Action blocked	Buttons disabled with message	Pass
9	GET /api/health on the deployed API	API and database status	status ok, database connected	Pass

Table 7.6 Automated Test Results (27 October 2026)

Suite	Command	Result
Back end (pytest, 12 files)	`python -m pytest -q`	170 passed
Front end (Vitest, 9 files)	`npx vitest run`	44 passed
Production build	`npm run build`	Succeeded

7.7 Implementation Screenshots

The following screenshots, captured from the deployed application, show the remaining implemented screens.

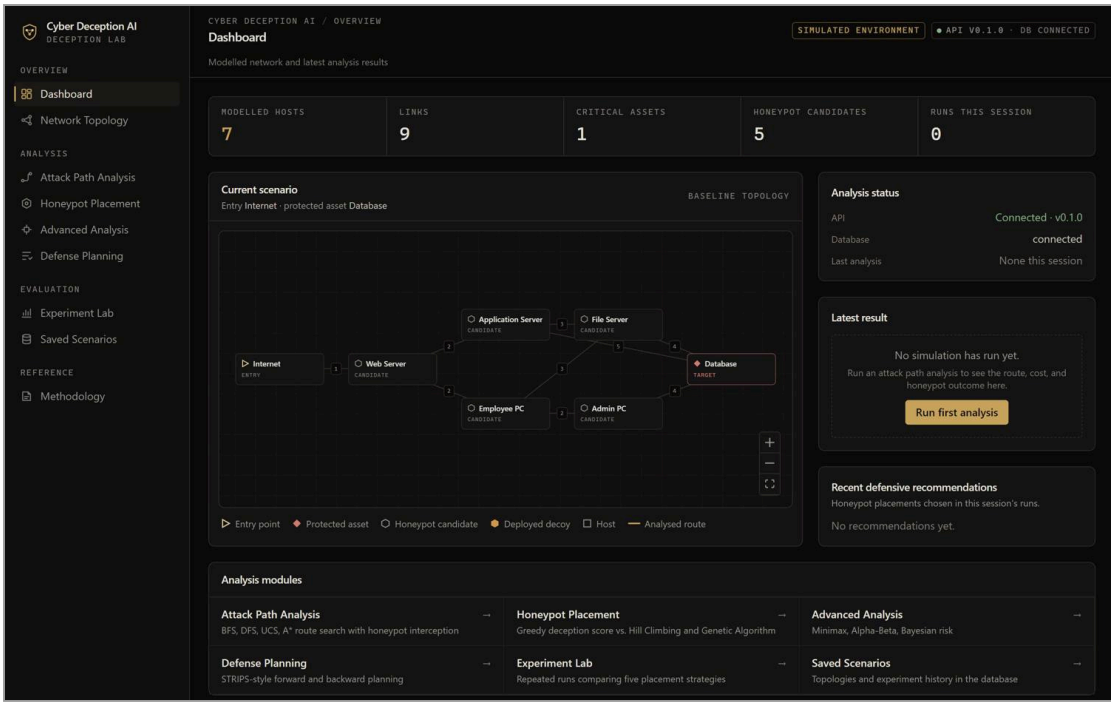


Fig. 7.7 Dashboard Before the First Analysis (Empty State)

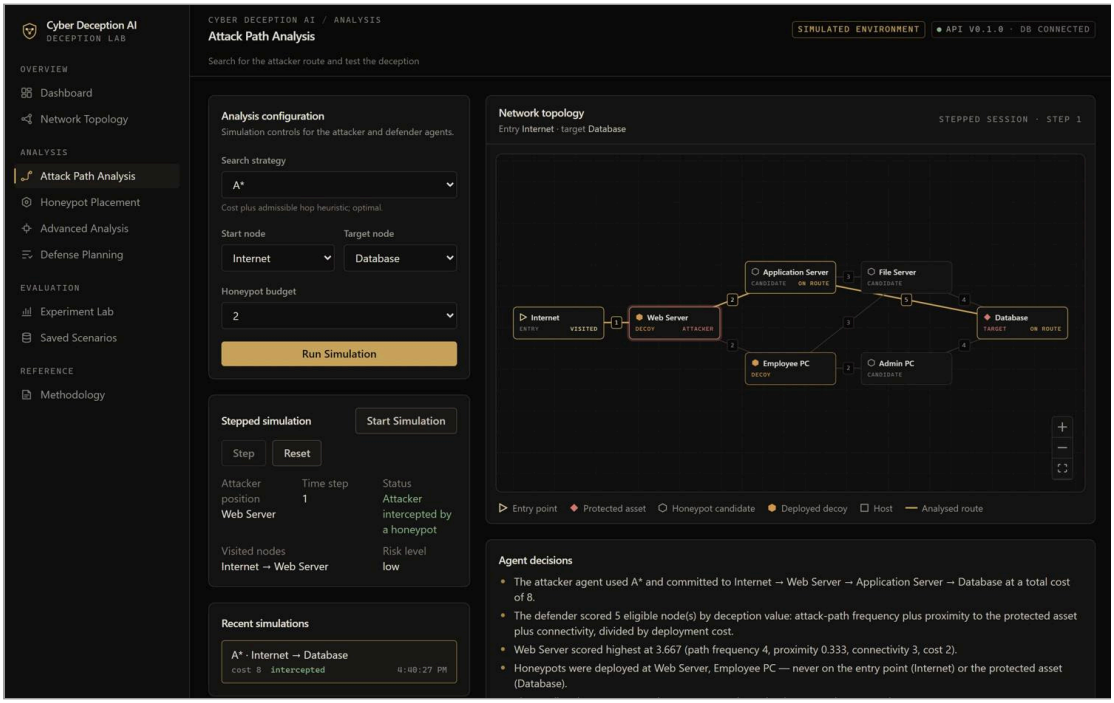


Fig. 7.8 Stepped Simulation: Attacker Intercepted at Web Server on Step 1

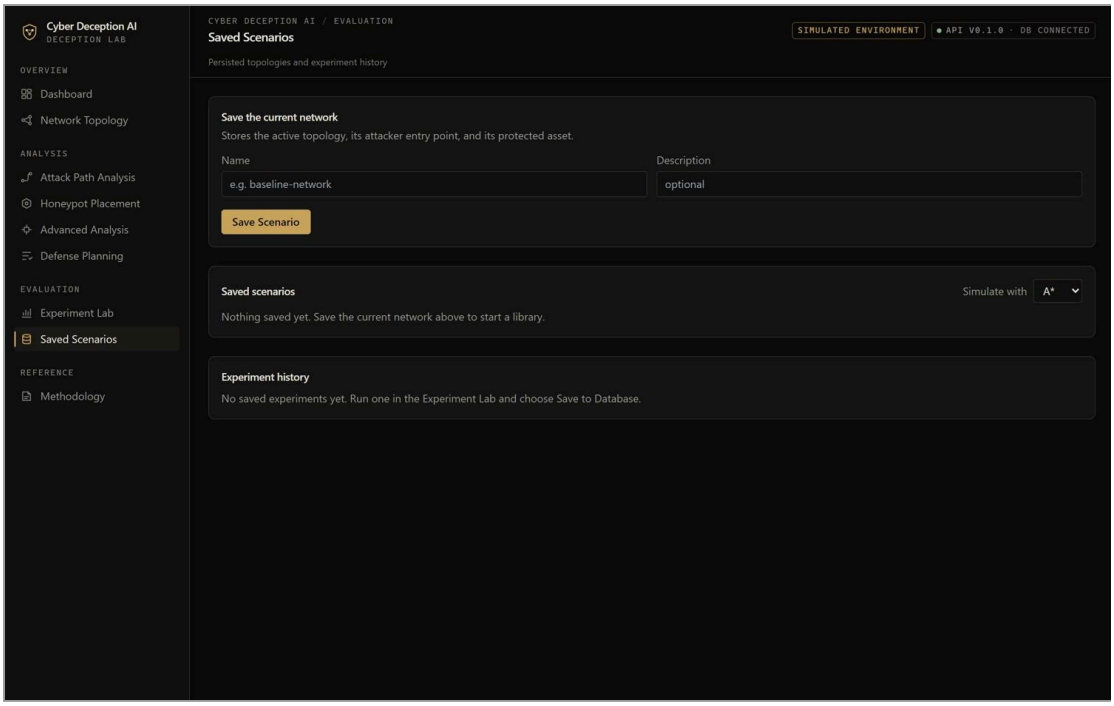


Fig. 7.9 Saved Scenarios Page Backed by the Database

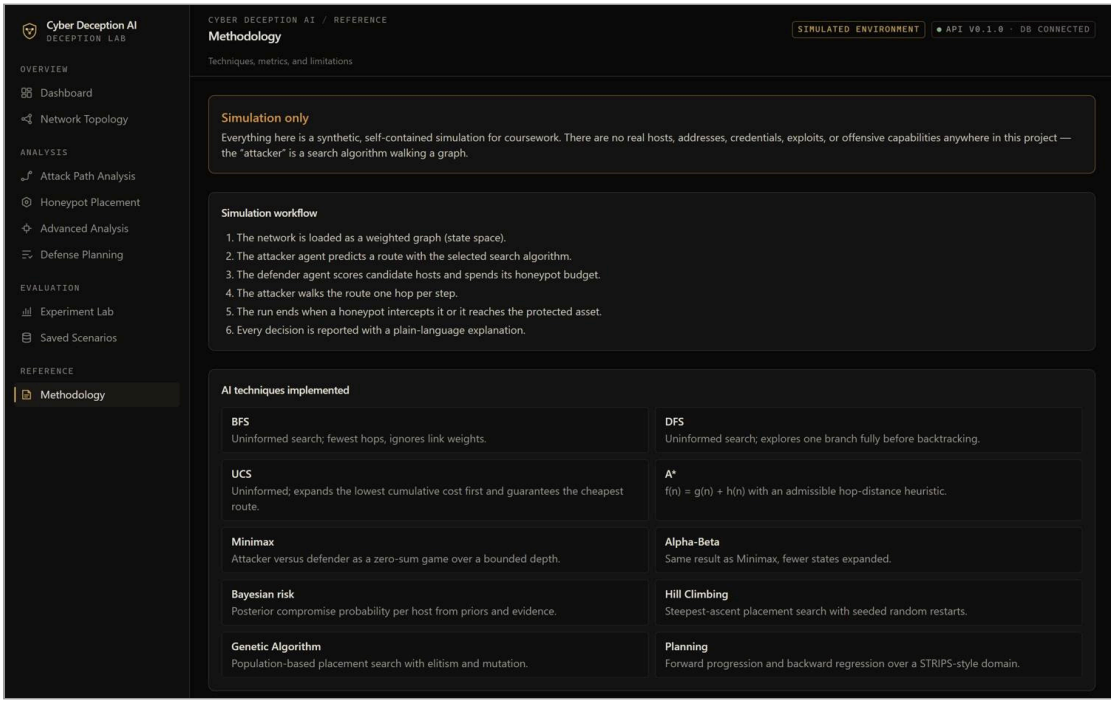


Fig. 7.10 Methodology Page

7.8 Key Source Code

The core loop of Uniform-Cost Search (`backend/ai/ucs.py`) pops the cheapest frontier entry and returns only when the goal is popped. This is what guarantees optimality:

```
while frontier:
    cost, _, current, path = heapq.heappop(frontier)
    if current in visited:
        continue
    visited.add(current)
    nodes_explored += 1
    if current == goal:
        return path, cost, nodes_explored
    for neighbor in graph.neighbors(current):
        edge_cost = graph[current][neighbor]["weight"]
        new_cost = cost + edge_cost
        if neighbor not in best_cost or new_cost <
best_cost[neighbor]:
            best_cost[neighbor] = new_cost
            counter += 1
            heapq.heappush(frontier, (new_cost, counter, neighbor,
path + [neighbor]))
return [], 0, nodes_explored
```

The greedy deception score (`backend/ai/optimizer.py`) implements Eqn. 6.1 and takes the top k candidates:

```
for node in candidates:
    freq = _attack_path_frequency(node, candidate_paths)
    proximity = round(_target_proximity(graph, node, target), 3)
    connectivity = _node_connectivity(graph, node)
    cost = (DEPLOYMENT_COST if deployment_costs is None else
deployment_costs).get(node, 1)
    raw_score = (freq + proximity + connectivity) / cost
    scored.append({"node": node, ..., "deception_score":
round(raw_score, 3)})
scored.sort(key=lambda x: x["deception_score"], reverse=True)
selected = [item["node"] for item in scored[:top_n]]
return scored, selected
```

7.9 Limitations

- A single seven-host synthetic topology with hand-set weights, costs and probabilities; all experiments are simulated.
- The attacker is assumed rational, cost-minimising and fully informed, and does not detect decoys.
- All four searches return the same route here, so the fitness route set contains one distinct route.
- The per-node greedy score ignores redundancy between decoys.
- Honeypot false positives and decoy isolation are not modelled.

8. CONCLUSION

This project addressed the problem of placing a limited number of honeypots in a network. It modelled the network as a weighted graph and predicted attacker routes with BFS, DFS, UCS and A*. It then compared three placement approaches on a common objective: a transparent Greedy score, Hill Climbing and a Genetic Algorithm. The work was delivered as a tested full-stack application with an interactive graph, explanations for every decision, an experiment laboratory and persistence. The application is publicly deployed on free hosting.

On the simulated network, all three algorithms reached the same fitness (6.2667) and covered all three cheapest routes with two honeypots. They intercepted the attacker in every run, compared with 50% (budget 1) and 85% (budget 2) for random placement, and Greedy was the fastest. These findings are limited by the network's chokepoint and by the simplified attacker model.

The main contribution is an explainable pipeline in which the defender's decision is computed directly from the attacker's predicted routes. Future work includes:

- larger, generated topologies without chokepoints;
- directed, vulnerability-aware graphs;
- dynamic attacker models;
- marginal-gain greedy and exact set-cover placement;
- game-theoretic defence;
- broader evaluation metrics with confidence intervals.

9. REFERENCES

1. L. Spitzner, *Honeypots: Tracking Hackers*. Boston, MA, USA: Addison-Wesley, 2003.
2. N. Provos, “A virtual honeypot framework,” in *Proc. 13th USENIX Security Symposium*, San Diego, CA, USA, 2004.
3. The MITRE Corporation, “MITRE Engage.” [Online]. Available: https://engage.mitre.org
4. J. Pawlick, E. Colbert, and Q. Zhu, “A game-theoretic taxonomy and survey of defensive deception for cybersecurity and privacy,” *ACM Computing Surveys*, vol. 52, no. 4, Art. 82, 2019, doi: 10.1145/3337772.
5. M. Tambe, *Security and Game Theory: Algorithms, Deployed Systems, Lessons Learned*. Cambridge, U.K.: Cambridge University Press, 2011.
6. D. E. Knuth and R. W. Moore, “An analysis of alpha-beta pruning,” *Artificial Intelligence*, vol. 6, no. 4, pp. 293–326, 1975.
7. O. Sheyner, J. Haines, S. Jha, R. Lippmann, and J. M. Wing, “Automated generation and analysis of attack graphs,” in *Proc. IEEE Symposium on Security and Privacy*, 2002, pp. 254–265, doi: 10.1109/SECPRI.2002.1004377.
8. X. Ou, S. Govindavajhala, and A. W. Appel, “MulVAL: A logic-based network security analyzer,” in *Proc. 14th USENIX Security Symposium*, Baltimore, MD, USA, 2005.
9. K. Durkota, V. Lisý, B. Bošanský, and C. Kiekintveld, “Optimal network security hardening using attack graph games,” in *Proc. 24th International Joint Conference on Artificial Intelligence (IJCAI)*, 2015.
10. E. W. Dijkstra, “A note on two problems in connexion with graphs,” *Numerische Mathematik*, vol. 1, pp. 269–271, 1959.
11. P. E. Hart, N. J. Nilsson, and B. Raphael, “A formal basis for the heuristic determination of minimum cost paths,” *IEEE Transactions on Systems Science and Cybernetics*, vol. 4, no. 2, pp. 100–107, 1968.
12. J. Y. Yen, “Finding the K shortest loopless paths in a network,” *Management Science*, vol. 17, no. 11, pp. 712–716, 1971, doi: 10.1287/mnsc.17.11.712.
13. A. A. Hagberg, D. A. Schult, and P. J. Swart, “Exploring network structure, dynamics, and function using NetworkX,” in *Proc. 7th Python in Science Conference (SciPy 2008)*, Pasadena, CA, USA, 2008, pp. 11–15.
14. T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, *Introduction to Algorithms*, 4th ed. Cambridge, MA, USA: MIT Press, 2022.
15. S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Hoboken, NJ, USA: Pearson, 2021.

16. G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher, “An analysis of approximations for maximizing submodular set functions—I,” *Mathematical Programming*, vol. 14, pp. 265–294, 1978.
17. J. H. Holland, *Adaptation in Natural and Artificial Systems*. Ann Arbor, MI, USA: University of Michigan Press, 1975.
18. D. E. Goldberg, *Genetic Algorithms in Search, Optimization, and Machine Learning*. Reading, MA, USA: Addison-Wesley, 1989.
19. R. E. Fikes and N. J. Nilsson, “STRIPS: A new approach to the application of theorem proving to problem solving,” *Artificial Intelligence*, vol. 2, no. 3–4, pp. 189–208, 1971.
20. FastAPI documentation. [Online]. Available: https://fastapi.tiangolo.com
21. React documentation. [Online]. Available: https://react.dev
22. Render documentation, “Blueprint specification.” [Online]. Available: https://render.com/docs/blueprint-spec
23. GitHub documentation, “Using custom workflows with GitHub Pages.” [Online]. Available: https://docs.github.com/en/pages

References [20]–[23] support the implementation and deployment descriptions in Sections 5 and 10.

10. DEPLOYMENT

The application is publicly deployed using only free services: GitHub Pages for the front end, and Render's free web-service and free PostgreSQL plans for the back end [22], [23]. Table 10.1 summarises the deployment, and Fig. 10.1 shows the live site as deployment evidence.

Table 10.1 Deployment Summary

Component	Service (plan)	Public URL	Verification
Front end	GitHub Pages (free)	`https://sid07-pro.github.io/cyber-deception-ai/`	Loads; API connected; simulation runs
Back-end API	Render web service (free)	`https://cyber-deception-api.onrender.com`	/api/health returns status ok
Database	Render PostgreSQL (free)	Internal to Render	Health reports database connected
Source code	GitHub	`https://github.com/sid07-pro/cyber-deception-ai`	Branch main

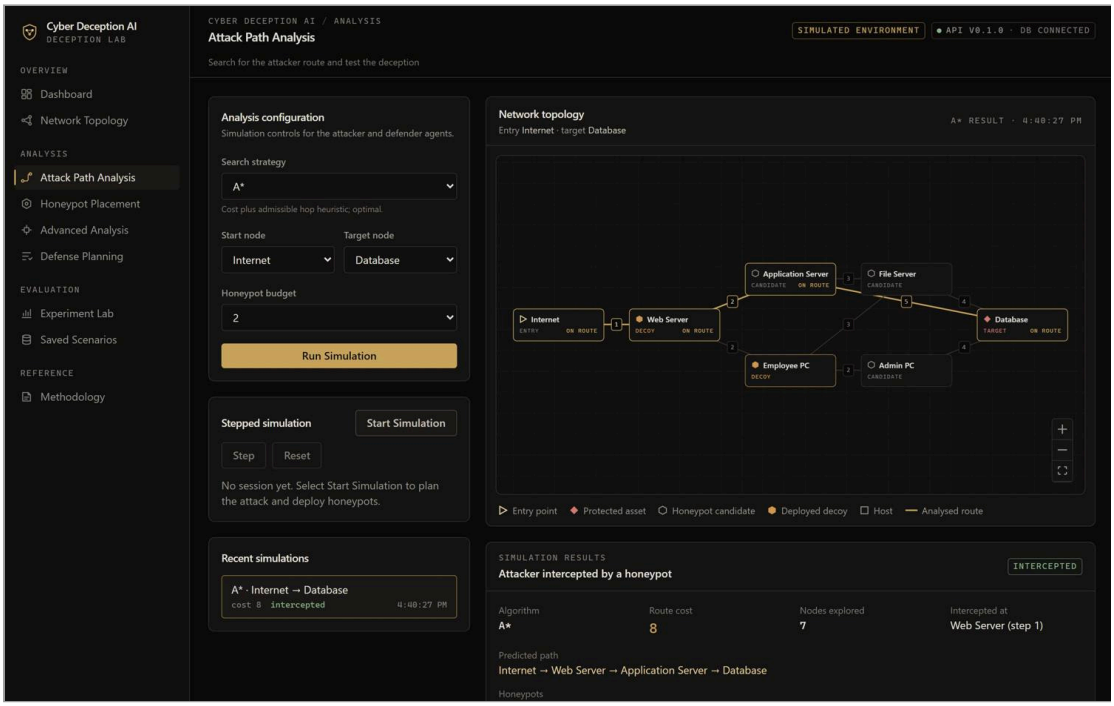


Fig. 10.1 Live Deployment on GitHub Pages

The live API health check returned the following response:

```
GET https://cyber-deception-api.onrender.com/api/health
{"status":"ok","service":"cyber-deception-ai","version":"0.1.0",
 "environment":"production","database":"connected"}
```

10.1 Deployment Process

- **Back end:** the repository's `render.yaml` blueprint creates the free web service (root `backend`, start command `uvicorn app.main:app --host 0.0.0.0 --port $PORT`, health check `/api/health`) and the free PostgreSQL database. Render injects `DATABASE_URL`.
- **CORS:** `ALLOWED_ORIGINS` is set to `https://sid07-pro.github.io`, so in production only the deployed front end may call the API.
- **Front end:** the GitHub Actions workflow `.github/workflows/pages.yml` builds the React app with `VITE_API_BASE_URL` set to the Render API and `VITE_BASE_PATH=/cyber-deception-ai/`, then publishes it to GitHub Pages on every push to `main`.

10.2 Local Execution

```
# Back end (Python 3.12)
cd backend
python -m venv .venv
.venv\Scripts\activate
pip install -r requirements.txt
uvicorn app.main:app --reload --port 8000

# Front end (Node.js 20+), in a new terminal
cd frontend
npm install
npm run dev            # http://localhost:5173

# Or the full stack with Docker Compose (set POSTGRES_PASSWORD in .env
first)
docker compose up -d --build
```

10.3 Testing

```
cd backend  && .venv\Scripts\python -m pytest -q
cd frontend && npx tsc -b --noEmit && npx vitest run && npm run build
```

10.4 Free-Tier Notes and Security

- No paid plan is used and no payment method is attached.
- The free Render web service sleeps after about 15 minutes of inactivity, so the first request afterwards can take up to a minute.
- Render's free PostgreSQL database is time-limited (about 30 days) and must then be recreated. Simulations do not depend on it.
- The system is simulation-only, with no real hosts, scanning or exploits. The API validates all input and restricts CORS, but it has no user authentication.

- No secrets are committed; database credentials are supplied by Render as environment variables.